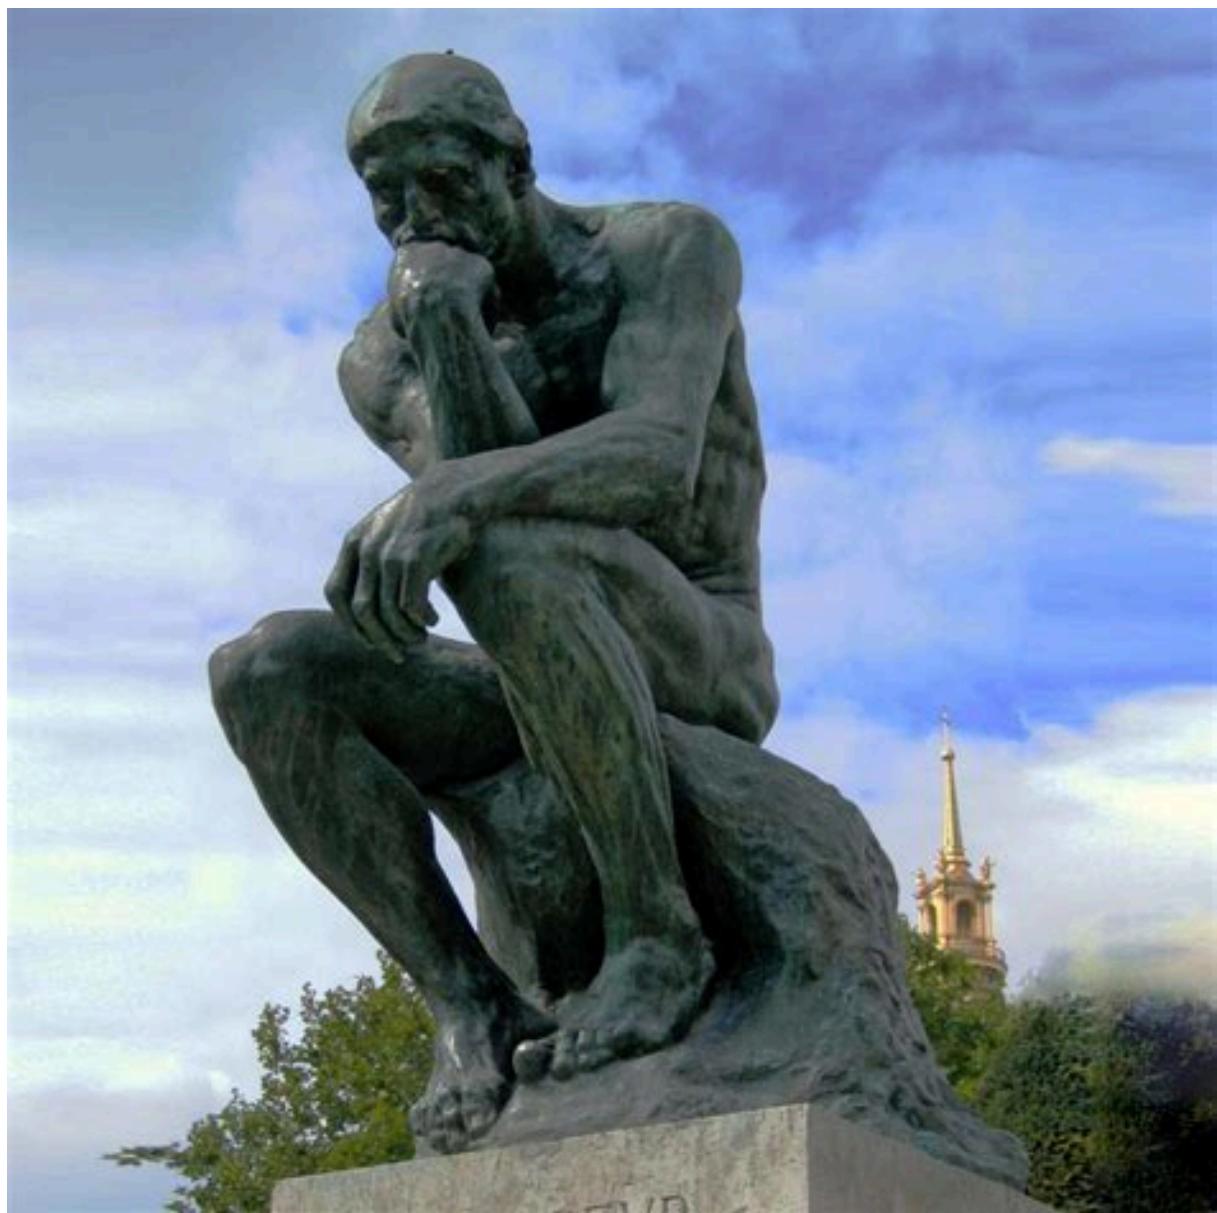

Large Language Models for Multi-Video Transcript Summarization

^aThorkild Kappel · ^bRahul TangellaPalli Sai Hanuma · ^cZhuolin Li

^{a,b,c}Aarhus University

^a202207326@post.au.dk · ^b202502250@post.au.dk · ^c202502085@post.au.dk



1 **1 Introduction and Motivation**

2 Information extraction and question answering are two of the key subjects when it comes to Natural
3 Language Processing (NLP) [1], [2], [3]. Ever since the new era of Artificial Intelligence, sparked by
4 the invention of the attention mechanism [4] as well as processing power to scale models to billions
5 of parameters, these subjects have changed in terms of methodology [5]. Summarization models are
6 typically trained using text to summarize datasets like TL;DR, which provide a training paradigm for
7 models to learn how to compress key information from lengthy texts [6], [7]. However, this fails to
8 address the scenario where the summary might benefit from multiple sources of information being
9 integrated into the summary together [8]. Using multiple sources instead could potentially mitigate
10 the risk of misinformation if sources disagree [9]. Further, having multiple perspectives and opinions
11 might reveal additional sides to view a problem from that a single source cannot show.

12 Video platforms such as YouTube represent an important medium for information seeking, especially
13 for explanatory queries [10]. YouTube transcripts are closer to spoken languages and typically less
14 structured, which often include extensive explanations, examples, and subjective opinions [11], and
15 are also typically more verbose [12]. In practice, users assess credibility and gather information
16 from multiple video-based sources when engaging with content on platforms like YouTube [13].
17 Therefore, in this project our goal is not to summarize a single transcript in isolation, but to make a
18 product that aggregates transcripts from multiple relevant videos based on user queries and generates
19 a comprehensive summary that answers the query while highlighting commonalities and differences
20 between videos when necessary.

21 Existing summarization datasets do not directly support tasks that rely on user queries or require
22 explicit comparison [14], [15]. To minimize this gap, we constructed a custom dataset. Target
23 summaries are generated through prompts to large language models, enabling subsequent fine-tuning
24 of smaller models. Through this process, the models learn to integrate information across multiple
25 transcripts and produce concise summaries. In our experiments, we compare a compact open-source
26 base model with two fine-tuned variants and evaluate their outputs against the reference summaries
27 using semantic similarity metrics.

28 **2 Methodology**

29 **2.1 Datasets and Preprocessing**

30 For this project, two datasets have been used. The first dataset used is the “Reddit TL;DR” dataset
31 [16], which has the purpose of being used for transfer learning, because of the large amount of data
32 it contains. A second custom dataset with less data has also been created, as the gold standard we
33 are trying to achieve.

34 The first dataset contains 116,722 posts scraped from Reddit that specifically contained a TL;DR
35 (Too Long Didn’t Read) section. This is a section in the post for readers who do not wish to read
36 the full post because it was too long. Once imported, we first split the data into samples that contain
37 for each post: the title, text in the post, and the TL;DR. Then downsampling was done to make each
38 subreddit have at max 2000 posts in our dataset. Afterwards, to put it into a similar format as the
39 second dataset for the purpose of better transfer learning, grouping of the posts based on similarity
40 was done using a greedy grouping algorithm.

41 To group posts together, first a TF-IDF vectorization (1) of the post texts was done. Then for each
42 post, their 50 most similar posts in terms of cosine similarity (4) were calculated. After this, starting
43 from one end of the data, formatted in no particular order, samples were grouped together with the
44 three posts that were the most similar, and that had not already been grouped in another group.
45 Furthermore, if samples are to be grouped, they must be above a cosine similarity of at least 0.3
46 and below 0.95 to ensure enough similarity and to ensure duplicate posts are not grouped together
47 respectively. For a group to be formed, it must also contain at least two samples, meaning that if all
48 50 closest samples are already in a group, we note the sample as ungrouped for now in the algorithm,
49 but it might be grouped at a later repetition. Groups that were ungrouped by the end were then simply
50 thrown away.

51 For concatenating a group into a single sample, we made a specific format. It is created by adding
52 each post one by one to a single string the delimited by a line separation, as well as prepending its
53 title and a post number between one and four. For grouping the TL;DR’s, 36 specific prepends were
54 made with the intention in mind that they should be similar to the second dataset. An example of
55 such a prepend that we used is “Transcript {a} discussed:”, where “{a}” will be replaced with the
56 transcript number, and the prepend is chosen at random for each TL;DR. Then for TL;DR’s in the
57 same group, a prepend was first attached, and the TL;DR’s were then concatenated into the same
58 text, delimited by a line space.

59 Secondly, another dataset was created by first generating 1109 prompts using Claude Sonnet 4.5 [17]
60 to some knowledge one might want to gain from watching a video. Examples of these include “how
61 to bake a pie”, “why is carbon monoxide deadly”, and “what is the sunk cost fallacy”. Then using
62 YouTube’s API [18] together with a YouTube transcript API [19]. These prompts were used to get the
63 top-four videos from YouTube and then extract their transcripts. Next these were put into a database
64 containing the prompt and video transcript pairs. Now, a single sample is constructed by grouping
65 together these four videos with the same YouTube prompt into one single text in the same format
66 as done with the TL;DR, except that they also have the YouTube prompt prepended. This means
67 samples will contain the YouTube prompt, and then for each video, the video number followed by
68 the title will be prepended on the transcript. Using these samples, outputs were generated by using
69 GPT 5.1 chat [20]. The instructions for the outputs were to create an output containing a separate
70 summary for each transcript that should cross-reference how they differ or agree with each other,
71 and should focus on answering the prompt using this information. The full prompt can be found in
72 the GitHub repository of the project (Appendix). This will be our gold standard example of how
73 the videos should be produced. Because of resource constraints due to large input sizes, inputs of
74 more than 10.000 tokens were filtered leading to 1004 samples. This was then split into a 90/10%
75 respectively train/test split leading to 903 train samples and 101 test samples.

76 **2.2 Model Fine-tuning Procedure**

77 For the purposes of comparison, two different models were fine-tuned. One model was fine-tuned
78 only on the second dataset (model 1), and another, both datasets (model 2). Both were based on
79 Llama 3.2, 3B, 8-bit quantized [21]. All fine-tuning was done in the same way using the same
80 hyper-parameters. For the purposes of staying within memory and time limits, LoRA fine-tuning
81 was utilized with around 24M trainable 16-bit parameters. LoRA adapters were injected into all
82 Attention and Forward Neural Network layers of the model (5) (6). The LoRA r parameter was set
83 to 16, meaning the LoRA downscaling parameter matrix (A) will have dimensions [3072 x 16] and
84 the upscaling matrix (B) will have size [16 x 3072], where 3072 is the model’s original embedding
85 dimensional size. The alpha scaling factor is set to 32, effectively scaling LoRA adapter weights by
86 $\frac{32}{16} = 2.0$. The learning rate was set to 2e-4. The amount of epochs was set to one for model 2 when
87 trained on the TL;DR dataset, and for both models, the amount of epochs was set to three for the
88 custom “gold standard” dataset, as this is the most important one.

89 **2.3 Model Testing**

90 For testing of the solution, we compared three models in terms of their output on the test set. The
91 first two models are explained above in section 2.2 (model 1 and model 2). The third model we test
92 is simply the standard 8-bit quantized LLama 3.2, 3B, used as a control model (model 0). To make
93 it fair, we gave the standard model the exact same instruction prepend that was given to GPT 5.1.
94 We then for each model gave it the input transcript from the test set, which neither of the models had
95 seen before, such that they generate the candidate results. Then we calculated the F1 BERTScore (9),
96 by comparing the candidate results and the GPT 5.1 outputs as the reference. This gave us 101 F1
97 BERTScores for each model. This overall process can be seen in Figure Figure 1. The embedding
98 model used in the BERTScore calculations was BART Large [22].

99 BERTScore is appropriate in our case, where there is no specifically correct answer. In comparison,
100 more “naive” word statistic methods like Bleu- or Rouge-1-scores will neglect the similarity of two
101 different words that have the same semantic meaning.

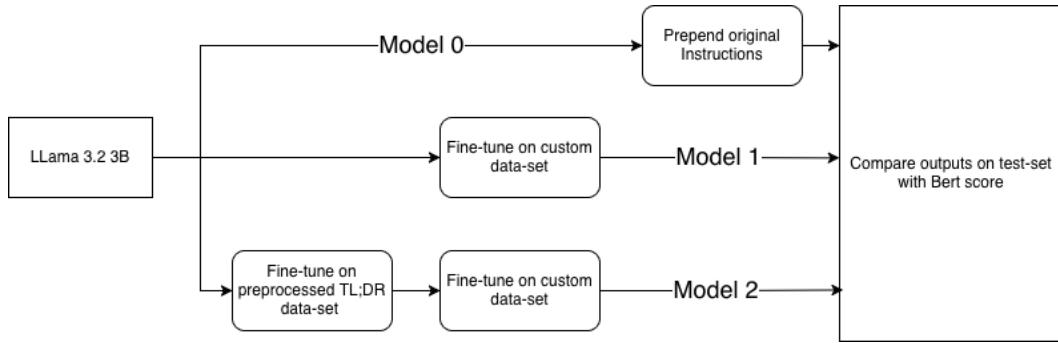


Figure 1: The testing workflow of the project

103 To compare the BERTScores of each, we conducted a Tukey's test [23] between the BERTScores of
104 the three groups.

105 **3 Results**

106 The results of Tukey's statistical test between the BERTScores of the models are shown in Table
107 Table 1 below:

Table 1: Results of the Tukey's statistical test

Comparison	statistic	p-value	Lower CI	Upper CI
(0 - 1)	-.053	< .001	-.062	-.043
(0 - 2)	-.039	< .001	-.049	-.030
(1 - 2)	.13	$p = .04$.004	.023

113 Further, the BERTScore distributions of the models can be seen plotted in the boxplots in the Figure
114 Figure 2 below:

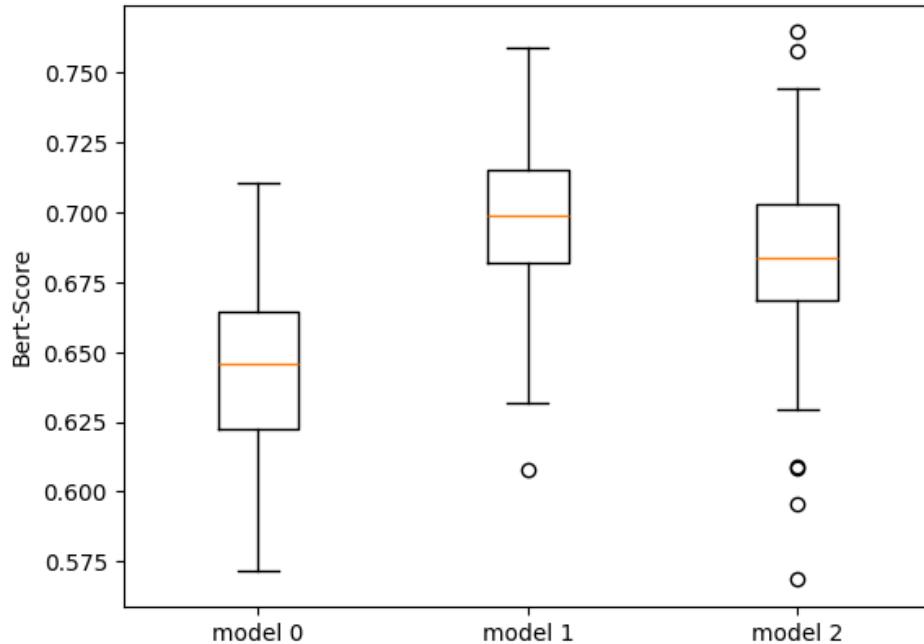


Figure 2: The testing workflow of the project

116 The results of the analysis show that, according to the BERTScore, both of the fine-tuned models
117 (models 1 and 2) performed significantly better than the non-fine-tuned base model, with both results
118 showing $p < .001$. However, looking at the Tukey statistic, we see that model 1, which was fine-
119 tuned only on the custom dataset, performs slightly better than model 2, with $p = .04$ and effect
120 .13. Inspecting the boxplots in Figure Figure 2 visually, taking outliers out of account, we can also
121 confirm the findings that model 1 performs better across the full distribution.

122 **4 Discussion**

123 **4.1 Summary and Discussion of the Results**

124 From the results, it is clear that the model fine-tuned only on our custom dataset performs the best.
125 This goes against the idea that using the TL;DR as an intermediate transfer learning dataset is useful.
126 The idea was that since the custom dataset contains a relatively small amount of data, we could
127 create a similar dataset from the TL;DR dataset. The TL;DR dataset has more data, so we could
128 use this to “get the model started”. However, this seems to have only introduced more noise for the
129 model, making it perform worse according to our results. We still see that it performs better than the
130 reference model, but this is expected since it is trained specifically on samples drawn from the same
131 dataset that we also tested the models on (although not the same samples). This drop in performance
132 of model 2 relative to model 1 could be because the two datasets are simply too dissimilar for it to
133 make sense. We expect that when trained on the reformatted TL;DR dataset that the model might have
134 lost some of its more general intelligence that the model originally had from pretraining. This also
135 seems to be in line with other results where overly fine-tuning language models actually made them
136 perform worse as they lose their pretrained features [24]. In terms of the BERTScore values in the
137 results, BERTScore does not generalize across tasks, and especially not across different embedding
138 models, so we only have the models in this project to compare with for now.

139 However, the result that model 1 performed the best is an important result, since for the actual
140 implementation of the full product pipeline, we now know that it will be better to use model 1.
141 Additionally, when inspecting some sample outputs, it was also noticeable that the base model did
142 not always create the expected format, whereas model 1 and model 2 were better at producing the
143 correct output format in accordance with the instructions.

144 Our project can thus be said to be a success in accordance with our goal of creating a product that
145 fetches multiple YouTube videos and summarizes and compares them with each other. Our project
146 provides a lightweight and computationally efficient way of doing this, in comparison to running
147 larger language models with larger sizes. Our model has around three billion parameters, which is a
148 small number in comparison to GPT 5.1’s amount of parameters. Although not publically disclosed,
149 more sources estimate that it contains more than one trillion parameters [25], [26] although, if it
150 is based on architectures like Mixture of Experts [27], it could mean that less parameters are used
151 for model inference. Thus, in comparison our model would save energy, time, and computational
152 resources, and open up the opportunity of running the model locally.

153 **4.2 Methodological Limitations**

154 Setting our reference to be GPT 5.1 generations has its limitations. Although GPT 5.1 is a high-
155 performing large language model as of this time, it is not perfect, and there is still room for
156 improvement [28]. Furthermore, we cannot inspect all the summaries that it generated. Although a
157 few sample outputs were inspected and were deemed to be well written, an expert human evaluation
158 could have potentially improved the reference summaries.

159 Important to mention is also that BertScore has its limitations in terms of evaluating performance.
160 Of course, since we are trying to generate a summary, there is no specific answer that is correct. This
161 in the first place, makes it hard to measure the performance of the model using quantitative means.
162 Qualitative methods, on the other hand, take time and effort to produce, especially given the 101
163 test examples that were made. BERTScore does not measure factual correctness, logical structure,
164 coherence, redundancy of text or human preferences (directly). It can only measure the similarity
165 of the contextual embeddings of each word with the most similar contextual embedding of a word
166 in the wanted output. And the “wanted output” in our case has not been fully evaluated. Although,

167 as written above, the few samples that were inspected worked well. In this way, human preference
168 could theoretically lean towards any of the three models, although the prompt was designed to align
169 with what we wanted. Thus, it has indirectly been aligned with human preference.

170 Since we used BART Large for the BERTScores, we also ran into a methodological problem. BART
171 Large has a maximum context size of 1024 tokens. This means that if we generated an output of
172 more than 1024 tokens that it would simply be truncated. While none of the GPT 5.1 outputs were
173 longer than that. It was found that seven of the outputs from model 0 were larger than this, meaning
174 they were automatically truncated by BART Large. However, these seven were mostly around 1100
175 tokens, meaning not many tokens were filtered out from these samples. This could potentially have
176 the effect of lowering the BERTScore slightly in these samples in comparison to what they could
177 have been, if some context in the GPT 5.1 output was missed from the filtered tokens. The same was
178 seen with model 1 where 13 generations were truncated, with the maximum amount of tokens in a
179 generation was 1054. Model 2, however, had 47 generations truncated, but the maximum amount of
180 tokens in a generation was only 1075 (see appendix for distributions). This could potentially lead
181 to an underestimation of its F1 BERTScores. However, since at max 51 tokens were filtered, this
182 effect might not have been huge and would probably still not have made it perform as well as model
183 1. Additionally, the cost of doing this extra fine-tuning in terms of computational resources, energy,
184 and time, is also a strong argument against doing it.

185 Lastly, sometimes YouTube transcripts do not contain a lot of talking. It could also be that the
186 transcripts fetched are simply not relevant to answering the questions. Although the summaries that
187 GPT 5.1 generated often took this into account, it is still not optimal. Especially if more videos
188 corresponding to the same prompt do not contain relevant information, this could lead to a lower
189 quality output.

190 **4.3 Avenues for Further Studies**

191 For further testing of our product, Comparison with larger language models on the specific task we
192 are trying to solve could be relevant. We have argued that this performs better on the same task
193 than similar sized standard models, but not that it performs better than general models with larger
194 parameter numbers. This will be necessary to understand when our model is useful, and how useful
195 it actually is.

196 Furthermore, we have not actually tested the full framework from start to end in one go. The next
197 steps would be to make a full implementation, where a user sends in a YouTube query, it goes into
198 the black box, and a summary of the top four videos is returned. This framework would have to be
199 further optimized for efficient run times and usage, as well as being tested for this. Also, human
200 testing would make preferential analysis much easier, but this is often costly to do.

201 Lastly, the project could be optimized by addressing the problem where transcripts do not contain
202 relevant information. A mechanism could be made to decide if transcripts are bad, and if they are,
203 other transcripts could potentially be fetched instead. Another way to address this is to make a more
204 intelligent searching mechanism that takes a look at the immediate available information for the
205 video, in order to decide if it is actually relevant for answering the question or not, in contrast to
206 simply taking the top four videos ordered by YouTube.

207 **5 Conclusion**

208 In summary, we have tested three models having the purpose of generating a summary of multiple
209 YouTube video transcripts, fetched from a specific search query. One model was a standard language
210 model. Another model was trained on a custom dataset created from YouTube transcripts that was
211 then summarized by GPT 5.1 as the target. The last model was trained using transfer learning from
212 first being trained on a similar but more abundant dataset before being fine-tuned on the custom
213 dataset. The test was conducted by comparing BERTScores between the models' output and a
214 specifically instructed generation from GPT 5.1, given the same transcripts. Our analysis showed
215 that the best performing model was the model that was only fine-tuned on the custom dataset. The
216 model showed promising results in terms of generating a summary of the videos, taking into account
217 factors such as referencing when video transcripts agreed and disagreed.

218 **Bibliography**

- 219 [1] R. Srihari and W. Li, "Information Extraction Supported Question Answering," Oct. 1999,
220 Accessed: Dec. 19, 2025. [Online]. Available: <https://apps.dtic.mil/sti/html/tr/ADA460042/>
- 221 [2] D. Moldovan and M. Surdeanu, "On the Role of Information Retrieval and Information
222 Extraction in Question Answering Systems," *Information Extraction in the Web Era*,
223 vol. 2700. Springer Berlin Heidelberg, Berlin, Heidelberg, pp. 129–147, 2003. doi:
224 10.1007/978-3-540-45092-4_6.
- 225 [3] O. Kolomiyets and M.-F. Moens, "A survey on question answering technology from an infor-
226 mation retrieval perspective," *Information Sciences*, vol. 181, no. 24, pp. 5412–5434, Dec.
227 2011, doi: 10.1016/j.ins.2011.07.047.
- 228 [4] A. Vaswani *et al.*, "Attention is All you Need," in *Advances in Neural Information Processing
229 Systems*, I. Guyon, U. V. Luxburg, S. Bengio, H. Wallach, R. Fergus, S. Vishwanathan, and R.
230 Garnett, Eds., Curran Associates, Inc., 2017. [Online]. Available: https://proceedings.neurips.cc/paper_files/paper/2017/file/3f5ee243547dee91fb053c1c4a845aa-Paper.pdf
- 232 [5] K. Kusunose, "Revolution of echocardiographic reporting: the new era of artificial intelligence
233 and natural language processing," *Journal of Echocardiography*, vol. 21, no. 3, pp. 99–104,
234 Sept. 2023, doi: 10.1007/s12574-023-00611-1.
- 235 [6] M. Völske, M. Potthast, S. Syed, and B. Stein, "Tl; dr: Mining reddit to learn automatic
236 summarization," in *Proceedings of the Workshop on New Frontiers in Summarization*, 2017,
237 pp. 59–63.
- 238 [7] S. Narayan, S. B. Cohen, and M. Lapata, "Ranking sentences for extractive summarization
239 with reinforcement learning," *arXiv preprint arXiv:1802.08636*, 2018.
- 240 [8] D. R. Radev, H. Jing, M. Styś, and D. Tam, "Centroid-based summarization of multiple
241 documents," *Information Processing & Management*, vol. 40, no. 6, pp. 919–938, 2004.
- 242 [9] E. Porter and T. J. Wood, "The global effectiveness of fact-checking: Evidence from simulta-
243 neous experiments in Argentina, Nigeria, South Africa, and the United Kingdom," *Proceedings
244 of the National Academy of Sciences*, vol. 118, no. 37, p. e2104235118, 2021, doi: 10.1073/
245 pnas.2104235118.
- 246 [10] M. Zimmermann and R. Jucks, "Investigating the Role of Communication for Information
247 Seekers' Trust-Related Evaluations of Health Videos on the Web," *Interactive Journal of
248 Medical Research*, vol. 7, no. 2, p. e10282, 2018, doi: 10.2196/10282.
- 249 [11] K. Kousha, M. Thelwall, and M. Abdoli, "The role of online videos in research communication:
250 A content analysis of Y ou T ube videos cited in academic publications," *Journal of the
251 American Society for information Science and Technology*, vol. 63, no. 9, pp. 1710–1727, 2012.
- 252 [12] D. Jurafsky and J. H. Martin, "Speech and Language Processing: An Introduction to Natural
253 Language Processing, Computational Linguistics, and Speech Recognition."
- 254 [13] E. M. Hughes, R. Wang, P. Juneja, T. W. Li, T. Mitra, and A. X. Zhang, "Viblio: Introducing
255 Credibility Signals and Citations to Video-Sharing Platforms," in *Proceedings of the CHI
256 Conference on Human Factors in Computing Systems*, in CHI '24. ACM, May 2024, pp. 1–20.
257 doi: 10.1145/3613904.3642490.
- 258 [14] Y. Xu and M. Lapata, "Generating Query Focused Summaries from Web Documents," in
259 *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics
260 (ACL)*, 2021. [Online]. Available: <https://aclanthology.org/2021.acl-long.475/>
- 261 [15] "Contrastive text summarization: a survey," *Complex & Intelligent Systems*, 2023, [Online].
262 Available: <https://link.springer.com/article/10.1007/s41060-023-00434-4>

- 263 [16] V. Cachola, K. Lo, A. Cohan, and D. S. Weld, “TL;DR: Extreme Summarization of Scien-
264 tific Documents,” in *Proceedings of the 2020 Conference on Empirical Methods in Natural*
265 *Language Processing (EMNLP)*, 2020.
- 266 [17] Anthropic, “Claude 4.5 Sonnet.” 2025.
- 267 [18] Google LLC, “YouTube Data API (v3).” 2024.
- 268 [19] “yt-dlp/yt-dlp.” Accessed: Dec. 19, 2025. [Online]. Available: <https://github.com/yt-dlp/yt-dlp>
- 269 [20] OpenAI, “ChatGPT (GPT-5.1).” 2025.
- 270 [21] Meta AI, “Llama 3.2: Open Foundation and Fine-Tuned Chat Models,” technical report, 2024.
271 [Online]. Available: <https://huggingface.co/meta-llama/Llama-3.2-3B>
- 272 [22] M. Lewis *et al.*, “BART: Denoising Sequence-to-Sequence Pre-training for Natural Language
273 Generation, Translation, and Comprehension,” in *Proceedings of the 58th Annual Meeting*
274 *of the Association for Computational Linguistics*, Association for Computational Linguistics,
275 2020, pp. 7871–7880. [Online]. Available: <https://aclanthology.org/2020.acl-main.703/>
- 276 [23] H. J. Keselman and J. C. Rogan, “The Tukey multiple comparison test: 1953–1976,” *Psycho-
277 logical Bulletin*, vol. 84, no. 5, pp. 1050–1056, 1977, doi: 10.1037/0033-2909.84.5.1050.
- 278 [24] A. Kumar, A. Raghunathan, R. Jones, T. Ma, and P. Liang, “Fine-Tuning can Distort Pretrained
279 Features and Underperform Out-of-Distribution.” Accessed: Dec. 19, 2025. [Online]. Available:
280 <http://arxiv.org/abs/2202.10054>
- 281 [25] “How Many Parameters does GPT-5 have - CometAPI - All AI Models in One API.” Accessed:
282 Dec. 19, 2025. [Online]. Available: [https://www.cometapi.com/en/how-many-parameters-
does-gpt-5-have/](https://www.cometapi.com/en/how-many-parameters-
283 does-gpt-5-have/)
- 284 [26] F. U. H. Zeya, “Comparing GPT-5.1, Gemini, and LLaMA 3: A Technical Deep
285 Dive.” Accessed: Dec. 19, 2025. [Online]. Available: [https://medium.com/@faizulhaquezeya/
comparing-gpt-5-1-gemini-and-llama-3-a-technical-deep-dive-0e58a5cd7c20](https://medium.com/@faizulhaquezeya/
286 comparing-gpt-5-1-gemini-and-llama-3-a-technical-deep-dive-0e58a5cd7c20)
- 287 [27] Y. Zhou *et al.*, “Mixture-of-Experts with Expert Choice Routing,” *Advances in Neural Infor-
288 mation Processing Systems*, vol. 35, pp. 7103–7114, Dec. 2022, Accessed: Dec. 19, 2025.
289 [Online]. Available: [https://proceedings.neurips.cc/paper_files/paper/2022/hash/2f00ecd787b
432c1d36f3de9800728eb-Abstract-Conference.html](https://proceedings.neurips.cc/paper_files/paper/2022/hash/2f00ecd787b
432c1d36f3de9800728eb-Abstract-Conference.html)
- 291 [28] M. D. McPhetridge, “GPT-5.1: A Structural Account of System Constraints.” Accessed: Dec.
292 19, 2025. [Online]. Available: <https://philpapers.org/rec/MCPGAS>

293 **6 Appendix**

294 **6.1 Equations**

295 TF-IDF (Term Frequency - Inverse Document Frequency):

$$\text{TF-IDF}(t, d) = \text{TF}(t, d) \times \text{IDF}(t) \quad (1)$$

296 where

$$\text{TF}(t, d) = \frac{f_{t,d}}{\sum_{t' \in d} f_{t',d}} \quad (2)$$

$$\text{IDF}(t) = \log\left(\frac{N}{1 + \text{DF}(t)}\right) \quad (3)$$

297 and $f_{t,d}$ is the frequency of term t in document d , N is the total number of documents in the corpus,
298 and $\text{DF}(t)$ is the document frequency (the number of documents containing term t at least once).

299 Cosine Similarity of two vectors:

$$\cos(\mathbf{u}, \mathbf{v}) = \frac{\mathbf{u} \cdot \mathbf{v}}{\|\mathbf{u}\| \|\mathbf{v}\|} \quad (4)$$

300 Normal linear layer

$$\mathbf{y} = \mathbf{W}\mathbf{x} \quad (5)$$

301 LoRA adapted linear layer

$$\mathbf{y} = \mathbf{W}\mathbf{x} + \frac{\alpha}{r} \mathbf{B}\mathbf{A}\mathbf{x} \quad (6)$$

302 Where the LoRA adaptation is scaled by α/r . In LoRA training, we let the weights of the original
303 \mathbf{W} matrix stay as is, and only train the downscaling Matrix \mathbf{A} and upscaling matrix \mathbf{B} by backprop-
304 agation and gradient descent as usual. \mathbf{x} is the input vector to be projected by the weights, e.g.
305 embeddings.

306 BERTScore Precision and Recall

$$\text{Precision} = \frac{1}{|\mathbf{C}|} \sum_{c \in \mathbf{C}} \max_{r \in \mathbf{R}} \cos(\mathbf{c}, \mathbf{r}) \quad (7)$$

$$\text{Recall} = \frac{1}{|\mathbf{R}|} \sum_{r \in \mathbf{R}} \max_{c \in \mathbf{C}} \cos(\mathbf{r}, \mathbf{c}) \quad (8)$$

307 Where \mathbf{C} is the matrix containing candidate contextual embeddings (The generations) and \mathbf{R} is the
308 matrix containing the reference contextual embeddings (The ground truth)

309 The F1 score is then harmonic mean of the precision and recall:

$$F_1 = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}} \quad (9)$$

310 **6.2 Figures**

311

distribution of tokens in outputs, with $n_tokens > 1024 = 7$

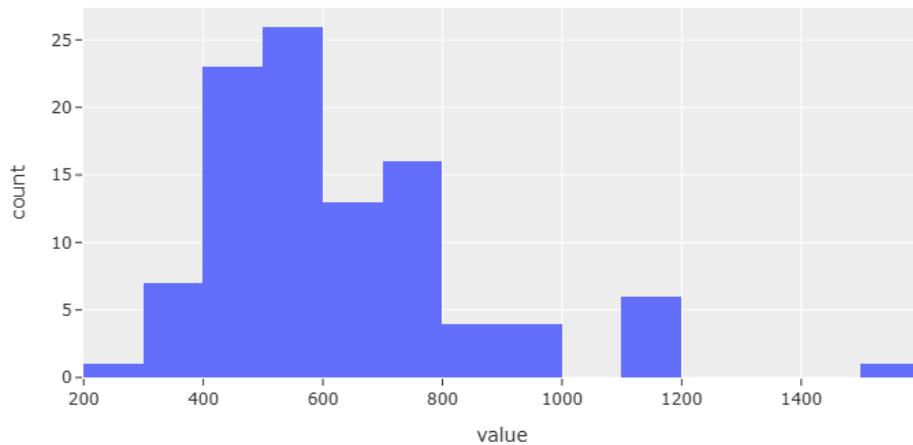
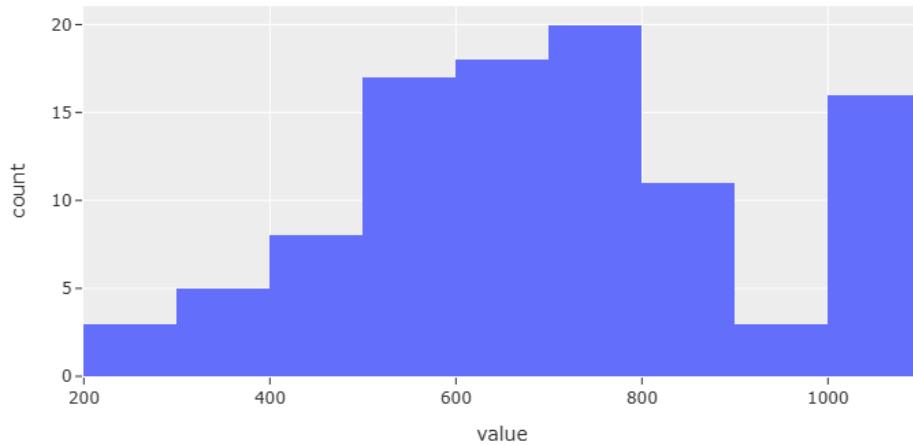


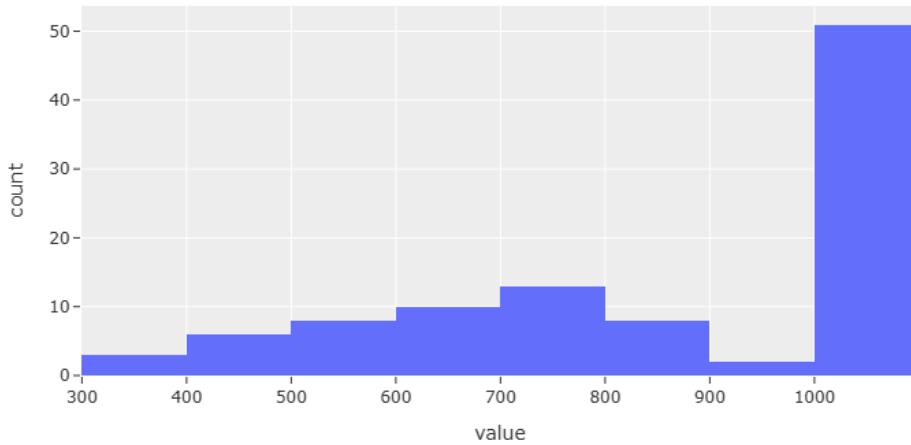
Figure 3: Output Tokens Length from Base Model (LLAMA 3.2 3B)

distribution of tokens in outputs, with $n_tokens > 1024 = 13$



312 Figure 4: Output Tokens Length from Fine-Tuned Model (LLAMA 3.2 3B with Custom Dataset)

distribution of tokens in outputs, with n_tokens > 1024 = 47



313 Figure 5: Output Tokens Length from Fine-Tuned Model (LLAMA 3.2 3B with TL;DR and Custom
314 Dataset)